

Milestone 2 - Submission Checklist

MLOps Course - Module 3

Pre-Submission Verification

Use this checklist to verify your Milestone 2 submission is complete before the deadline. Work through each section systematically.

Deliverables Checklist

Part 1: Multi-Stage Docker Image (4 points)

- Dockerfile exists with multi-stage build (builder + runtime stages)
- app/ directory contains inference script (e.g., app.py)
- app/requirements.txt or app/pyproject.toml with pinned dependencies
- docker-compose.yaml present (optional but recommended)
- Registry verification screenshot or link saved/documentated

Part 2: CI/CD Pipeline (4 points)

- .github/workflows/build.yml exists and is valid YAML
- Workflow includes test job with pytest
- Workflow includes build job for Docker image
- Workflow includes registry authentication step
- Workflow includes publish step with semantic version tags
- README.md contains CI/CD status badge
- README.md includes image pull/run instructions
- tests/test_app.py exists with unit tests

Part 3: Operations Runbook (2 points)

- RUNBOOK.md exists with all required sections:
 - Dependency pinning strategy documented
 - Image optimization with before/after size metrics
 - Security considerations explained
 - CI/CD workflow step-by-step explanation
 - Versioning strategy (semantic versioning)
 - Troubleshooting guide with common issues

Submission Requirements

-
- All files pushed to `module3/milestone2/` directory
 - Final commit tagged as `m2-submission`
 - CI/CD pipeline shows green (passing) status
 - Image accessible in course container registry
-

Rubric Evidence Map

Use this table to verify you have evidence for each graded criterion:

Criteria	Points	Evidence Location	Verification
Multi-stage Build	2	Dockerfile lines showing <code>FROM ... as builder</code> and second <code>FROM</code>	[] Verified
Runtime Environment	1	<code>requirements.txt</code> with pinned versions (e.g., <code>flask==2.3.2</code>)	[] Verified
Registry Integration	1	Screenshot or registry URL with <code>v1.0.0</code> style tag	[] Verified
Pipeline Functionality	2	<code>.github/workflows/build.yml</code> [] Verified with test→build→push jobs	
Test Integration	1	<code>tests/test_app.py</code> + workflow showing <code>pytest</code> step	[] Verified
Authentication & Versioning	1	Workflow secrets usage + semantic version in image tag	[] Verified
Runbook Quality	1	<code>RUNBOOK.md</code> with all 6 required sections	[] Verified
Project Organization	0.5	Clean directory structure under <code>module3/milestone2/</code>	[] Verified
README & Instructions	0.5	<code>README.md</code> with badge + quick start guide	[] Verified

Common Pitfalls

** Docker Optimization Pitfalls**

- Installing dev dependencies in runtime image (use multi-stage to avoid this)
- Not ordering Dockerfile commands from least to most frequently changing

-
- Using full base images instead of `-slim` or Alpine variants
 - Missing `.dockerignore` causing bloated images

** CI/CD Pitfalls**

- Forgetting to authenticate before pushing to registry
- Tests passing locally but failing in CI due to environment differences
- Hardcoded credentials instead of using GitHub Secrets
- Push step running even when tests fail (missing needs: test)

** Documentation Pitfalls**

- Outdated documentation that doesn't match current implementation
 - Missing exact commands (only descriptions)
 - Assuming reader familiarity with your setup
 - No CI badge in README
-

Automated Sanity Checks

Run these commands from your `module3/milestone2/` directory before submitting:

File Existence Checks

```
# Check all required files exist
echo "== Checking required files =="
for file in Dockerfile README.md RUNBOOK.md .github/workflows/build.yml; do
    if [ -f "$file" ]; then
        echo "/ $file exists"
    else
        echo "x MISSING: $file"
    fi
done

# Check app directory
if [ -d "app" ]; then
    echo "/ app/ directory exists"
    ls -la app/
else
    echo "x MISSING: app/ directory"
fi

# Check tests directory
if [ -d "tests" ]; then
    echo "/ tests/ directory exists"
    ls -la tests/
else
    echo "x MISSING: tests/ directory"
fi
```

Dockerfile Validation

```
# Verify multi-stage build structure
echo "==== Checking Dockerfile structure ==="
if grep -q "FROM.*as builder" Dockerfile 2>/dev/null || grep -q "FROM.*AS
    builder" Dockerfile 2>/dev/null; then
    echo "✓ Builder stage found"
else
    echo "✗ Missing builder stage (check 'FROM ... as builder')"
fi

# Count FROM statements (should be at least 2 for multi-stage)
FROM_COUNT=$(grep -c "^FROM" Dockerfile 2>/dev/null || echo 0)
if [ "$FROM_COUNT" -ge 2 ]; then
    echo "✓ Multi-stage build detected ($FROM_COUNT stages)"
else
    echo "✗ Not a multi-stage build (only $FROM_COUNT FROM statement)"
fi
```

Dependency Pinning Check

```
# Check for pinned dependencies
echo "==== Checking dependency pinning ==="
if [ -f "app/requirements.txt" ]; then
    UNPINNED=$(grep -E "^[a-zA-Z]" app/requirements.txt | grep -v "==" | grep
        -v ">=" | head -5)
    if [ -z "$UNPINNED" ]; then
        echo "✓ Dependencies appear pinned"
    else
        echo " Potentially unpinned dependencies found:"
        echo "$UNPINNED"
    fi
elif [ -f "app/pyproject.toml" ]; then
    echo "✗ Using pyproject.toml - manually verify pinning"
else
    echo "✗ No requirements.txt or pyproject.toml in app/"
fi
```

Docker Build Test

```
# Test Docker build locally
echo "==== Testing Docker build ==="
docker build -t milestone2-test:local . && echo "✓ Docker build successful"
|| echo "✗ Docker build failed"
```

```
# Check image size
docker images milestone2-test:local --format "Image size: {{.Size}}"
```

Test Execution

```
# Run tests locally
echo "== Running tests =="
pip install pytest -q
pytest tests/ -v && echo "✓ All tests passed" || echo "✗ Tests failed"
```

GitHub Actions Workflow Validation

```
# Validate workflow YAML syntax
echo "== Validating workflow YAML =="
if command -v python3 &> /dev/null; then
    python3 -c "import yaml;
        yaml.safe_load(open('.github/workflows/build.yml'))" 2>/dev/null && \
    echo "✓ Workflow YAML is valid" || echo "✗ Workflow YAML has syntax
        errors"
else
    echo "✗ Python not available - manually check YAML syntax"
fi

# Check for required workflow components
echo "== Checking workflow components =="
WORKFLOW=".github/workflows/build.yml"
grep -q "pytest" "$WORKFLOW" && echo "✓ pytest step found" || echo "✗
    Missing pytest step"
grep -q "docker.*build" "$WORKFLOW" && echo "✓ Docker build step found" || \
    echo "✗ Missing Docker build"
grep -q "docker.*push" "$WORKFLOW" && echo "✓ Docker push step found" || \
    echo "✗ Missing Docker push"
grep -q "secrets\." "$WORKFLOW" && echo "✓ Secrets usage found" || echo "✗
    No secrets referenced"
```

README Validation

```
# Check README contains required elements
echo "== Checking README =="
if grep -q "badge" README.md 2>/dev/null || grep -q "\[" README.md
    2>/dev/null; then
    echo "✓ Badge or image found in README"
else
    echo "✗ No CI badge found in README"
fi
```

```
if grep -q "docker pull\|docker run" README.md 2>/dev/null; then
    echo "✓ Docker pull/run instructions found"
else
    echo "✗ Missing docker pull/run instructions"
fi
```

Git Tag Check

```
# Check for submission tag
echo "== Checking Git tags =="
if git tag | grep -q "m2-submission"; then
    echo "✓ m2-submission tag exists"
else
    echo "✗ Missing m2-submission tag (run: git tag m2-submission && git push
        --tags)"
fi
```

Self-Assessment Questions

Before submitting, honestly answer these questions:

Reproducibility

- Can someone clone my repo and build the Docker image in under 5 minutes?**
- Have I tested building from a clean clone (not just my local cache)?
- Are all dependencies pinned to specific versions?

CI/CD Pipeline

- Does my CI badge show green on the main branch?**
- If I push a commit with failing tests, does the pipeline correctly fail?
- Does the workflow only push images after tests pass?

Docker Quality

- Is my final image size reasonable (<500MB for most ML services)?
- Have I excluded unnecessary files using `.dockerignore`?
- Does my container run as a non-root user?

Documentation

- Could a team member understand my system from `RUNBOOK.md` alone?
- Are all commands in my documentation copy-pasteable?
- Have I documented what to do when common issues occur?

Security

-
- Are there any hardcoded credentials in my repository?
 - Have I used GitHub Secrets for registry authentication?
 - Does my Docker image use a minimal attack surface?

Versioning

- Does my registry image use semantic versioning (e.g., v1.0.0)?
 - Is my submission tagged as m2-submission in Git?
-

Quick Reference: Expected Repository Structure

```
module3/milestone2/
├── .github/
│   └── workflows/
│       └── build.yml           # CI/CD pipeline
├── app/
│   ├── app.py                 # Main inference script
│   └── requirements.txt       # Pinned dependencies
├── tests/
│   └── test_app.py            # Unit tests
├── .dockerignore             # Files to exclude from build
├── Dockerfile                # Multi-stage build
├── docker-compose.yaml        # (Optional) Local development
├── README.md                 # Project docs with CI badge
└── RUNBOOK.md                # Operations documentation
```

Final Submission Checklist

Complete this final checklist immediately before submitting:

- All automated sanity checks pass
- All self-assessment questions answered “yes”
- CI/CD pipeline shows green status
- m2-submission tag pushed to remote
- Image accessible in course registry with semantic version tag
- Reviewed rubric evidence map - all criteria verified